

# Lending Club - Loan Repayment Prediction Model - Final Report

## Introduction

The goal of each and every business is to make profit. For a lender, profit depends on whether or not the borrower repays the principal as well as interest. Without repayment the lender will incur a loss and that loss can even potentially be greater than the initial loan amount when lawyer, court and collection fees are taken into consideration. Therefore, it is critically important for a lender to be able to identify whether a potential borrower can and will make all of his or her loan payments. The purpose of this project is to identify the characteristics (limited to those found in the dataset) of persons who are likely to default on their loans and provide a simple framework for borrowers to make such a distinction. In order to make predictions on which potential borrower is or isn't likely to default, binary classification predictive models will be created using a variety of machine learning algorithms. The framework for lenders to make predictions concerning potential borrowers will come in the form of a decision tree.

## Literature Review

### Benchmarking state-of-the-art classification algorithms for credit scoring: A ten-year update<sup>2</sup>

[http://www.business-school.ed.ac.uk/waf/crc\\_archive/2013/42.pdf](http://www.business-school.ed.ac.uk/waf/crc_archive/2013/42.pdf) ([http://www.business-school.ed.ac.uk/waf/crc\\_archive/2013/42.pdf](http://www.business-school.ed.ac.uk/waf/crc_archive/2013/42.pdf))

Essentially this paper is a modern update to the landmark benchmarking study of classification algorithms for credit scoring by Baesens, Van Gestel, Viaene, Stepanova, Suykens and Vanthienen in 2003<sup>1</sup>. Since 2003 many new techniques and algorithms have been developed in predictive modeling<sup>2</sup>. This paper builds upon Baesens et al. by including all of the newer state-of-the-art techniques as well as those covered in the previous study. I chose this paper because its goal of comparing of classification algorithms for credit scoring is highly related to the problem this project aims to solve. The paper describes a vector  $x$  of  $m$  dimensions with each dimension as a feature characterizing an application for a credit products such as a loan<sup>2</sup>. The dataset used in this project is in the same format, where each row (vector) represents an individual loan. The study then goes on to discuss a binary response variable which indicates the existence or non-existence of a default event<sup>2</sup>. The probability of a default event given  $x$  is the classification problem being addressed in the study. Finally, a decision maker will take this probability and if it falls under a given threshold the application will be accepted, otherwise it will be rejected. This is essentially the approach I am taking for this capstone project. The Lending Club data set I am using for this project contains over 100 features characterizing the borrower, my vector  $x$ . The data set also contains a feature stating the current status of the loan with various possible values, each of which can easily be grouped into default or non-default statuses. This is effectively my binary response variable as described in the study. Finally my goal is to estimate the probability of default given a set of borrower characteristics and use that to determine whether they are likely or unlikely to default. Again, this is parallel to the study. The study considered the following classification algorithms.

TABLE 2: CLASSIFICATION ALGORITHMS CONSIDERED IN THE BENCHMARKING STUDY

	Base model selection	Classification algorithm	Acronym	Number of models <sup>1</sup>
Individual classifier	n.a.	Bayesian Network	B-Net	4
		CART	CART	10
		Extreme learning machine	ELM	120
		Kernalized ELM	ELM-K	200
		k-nearest neighbor	kNN	22
		J4.8	J4.8	36
		Linear discriminant analysis <sup>2</sup>	LDA	1
		Linear support vector machine	SVM-L	29
		Logistic regression <sup>2</sup>	LR	1
		Multilayer perceptron artificial neural network	ANN	171

	Naive Bayes	NB	1		
	Quadratic discriminant analysis <sup>2</sup>	QDA	1		
	Radial basis function neural network	RbfNN	5		
	Regularized logistic regression	LR-R	27		
	SVM with radial basis kernel function	SVM- Rbf	300		
	Voted perceptron	VP	5		
Classification models from individual classifiers		16	933		
Homogenous ensembles	n.a.	Alternating decision tree	ADT	5	
		Bagged decision trees	Bag	9	
		Bagged MLP	BagNN	4	
		Boosted decision trees	Boost	48	
		Logistic model tree	LMT	1	
		Random forest	RF	30	
		Rotation forest	RotFor	25	
		Stochastic gradient boosting	SGB	9	
Classification models from homogeneous ensembles		8	131		
Heterogeneous ensembles	n.a.	Simple average ensemble	AvgS	1	
		Weighted average ensemble	AvgW	1	
	Static direct		Complementary measure	CompM	4
			Ensemble pruning via reinforcement learning	EPVRL	4
			GASEN	GASEN	4
			Hill-climbing ensemble selection	HCES	12
			HCES with bootstrap sampling	HCES-Bag	16
			Matching pursuit optimization of ensemble classifiers	MPOCE	1
			Stacking	Stack	6
			Top- $T$ ensemble	Top- $T$	12
	Static indirect		Clustering using compound error	CuCE	1
			k-Means clustering	k-Means	1
			Kappa pruning	KaPru	4
			Margin distance minimization	MDM	4
			Uncertainty weighted accuracy	UWA	4
	Dynamic		Probabilistic model for classifier competence	PMCC	1
			k-nearest oracle	kNORA	1
Classification models from heterogeneous ensembles		17	77		
Overall number of classification algorithms and models		41	1141		

The following table measures the performance of each algorithm in credit scoring classification using Area Under a ROC Curve (AUC). According to the study across all performance measures the top three most accurate classifiers are Random Forests, Bagged (MLP) Neural Networks, and Bagged Decision Trees.

TABLE 5: PERFORMANCE OF INDIVIDUAL CLASSIFIERS AND HOMOGENEOUS ENSEMBLES IN TERMS OF THE AUC

		AC	GC	Bene1	Bene2	UK	PAK	GMC
Individual classifiers	ANN	.926 (.011)	.791 (.014)	<u>.791</u> (.009)	<u>.802</u> (.005)	<u>.742</u> (.008)	<u>.644</u> (.004)	.859 (.003)
	B-Net	.922 (.011)	.764 (.015)	.771 (.009)	.786 (.009)	.703 (.023)	.623 (.004)	<u>.860</u> (.003)
	CART	.856 (.019)	.706 (.031)	.706 (.021)	.713 (.021)	.684 (.012)	.565 (.015)	.797 (.025)
	ELM	.911 (.011)	.778 (.012)	.766 (.010)	.761 (.006)	.650 (.009)	.599 (.003)	.717 (.004)
	ELM-K	.926 (.012)	.794 (.015)	.787 (.007)	.788 (.005)	.734 (.009)	.643 (.004)	.702 (.004)
	J4.8	.915 (.014)	.734 (.020)	.761 (.012)	.747 (.011)	.500 (.000)	.500 (.000)	.500 (.000)
	k-NN	.906 (.016)	.772 (.010)	.765 (.009)	.754 (.007)	.725 (.014)	.600 (.005)	.739 (.004)
	LDA	.929 (.009)	.784 (.012)	.775 (.011)	.779 (.008)	.715 (.010)	.626 (.003)	.692 (.004)
	LR	<u>.931</u> (.011)	.784 (.012)	.773 (.012)	.791 (.006)	.720 (.011)	.626 (.003)	.693 (.005)
	LR-R	.925 (.012)	.778 (.015)	.787 (.007)	.798 (.004)	.690 (.012)	.635 (.004)	.623 (.006)
	NB	.893 (.020)	.777 (.017)	.747 (.013)	.724 (.010)	.701 (.019)	.613 (.006)	.671 (.003)
	RbfNN	.902 (.019)	.762 (.013)	.760 (.009)	.739 (.007)	.701 (.014)	.604 (.003)	.755 (.007)
	QDA	.917 (.018)	.674 (.148)	.765 (.011)	.780 (.006)	.703 (.012)	.612 (.004)	.811 (.003)
	SVM-L	.924 (.013)	.782 (.014)	.786 (.007)	.796 (.003)	.659 (.014)	.636 (.004)	.733 (.017)
	SVM-Rbf	.926 (.012)	<u>.799</u> (.011)	.786 (.008)	.795 (.004)	.666 (.028)	.630 (.004)	.815 (.009)
	VP	.810 (.030)	.680 (.020)	.698 (.013)	.621 (.017)	.554 (.018)	.567 (.003)	.568 (.024)
Homogeneous ensemble classifiers	ADT	.929 (.010)	.758 (.012)	.786 (.008)	.794 (.010)	.732 (.008)	.641 (.004)	.860 (.004)
	Bag	.930 (.014)	.788 (.014)	.794 (.008)	.805 (.006)	.742 (.007)	.643 (.003)	<u>.864</u> (.003)
	BagNN	.927 (.012)	<u>.802</u> (.010)	.793 (.008)	.802 (.004)	<u>.745</u> (.008)	<u>.646</u> (.004)	.838 (.004)
	Boost	.930 (.010)	.772 (.012)	<u>.795</u> (.007)	<u>.808</u> (.005)	.741 (.010)	.643 (.004)	.860 (.003)
	LMT	.930 (.013)	.747 (.015)	.780 (.007)	.787 (.006)	.720 (.010)	.630 (.004)	.833 (.017)
	RF	<u>.931</u> (.014)	.789 (.013)	.794 (.008)	.805 (.006)	.742 (.007)	.643 (.003)	<u>.864</u> (.003)
	RotFor	.929 (.013)	.773 (.015)	.788 (.007)	.794 (.007)	.502 (.016)	.635 (.002)	.820 (.005)
	SGB	.928 (.013)	.751 (.015)	.786 (.007)	.797 (.006)	.735 (.012)	.642 (.004)	.860 (.003)

### An Empirical Comparison of Supervised Learning Algorithms<sup>3</sup>

<https://www.cs.cornell.edu/~caruana/ctp/ct.papers/caruana.icml06.pdf>

(<https://www.cs.cornell.edu/~caruana/ctp/ct.papers/caruana.icml06.pdf>)

This study compares the performance of eight machine learning algorithms namely, SVMs, neural nets, logistic regression, naïve bayes, memory based learning, random forests, decision trees, bagged trees, boosted trees, and boosted stumps. The performance metrics used are, accuracy, F-score, Lift, ROC Area, average precision, squared error and cross entropy. The study concludes that bagged trees, random forests and neural nets have the best average performance (prior to calibration) over all the metrics and over all the problems. When calibration is taken into account, the overall best performing algorithm is boosted decision trees (calibrated)<sup>3</sup>. In close second is Random forests, followed by bagged decision trees (uncalibrated).

Table 3. Normalized scores of each learning algorithm by problem (averaged over eight metrics)

MODEL	CAL	COVT	ADULT	LTR.P1	LTR.P2	MEDIS	SLAC	HS	MG	CALHOUS	COD	BACT	MEAN
BST-DT	PLT	<b>.938</b>	.857	<b>.959</b>	<b>.976</b>	.700	.869	<b>.933</b>	.855	<b>.974</b>	<b>.915</b>	.878*	<b>.896*</b>
RF	PLT	.876	.930	.897	.941	<b>.810</b>	.907*	.884	.883	.937	.903*	.847	.892
BAG-DT	—	.878	.944*	.883	.911	.762	.898*	.856	<b>.898</b>	.948	.856	<b>.926</b>	.887*
BST-DT	ISO	.922*	.865	.901*	.969	.692*	.878	.927	.845	.965	.912*	.861	.885*
RF	—	.876	.946*	.883	.922	.785	.912*	.871	.891*	.941	.874	.824	.884
BAG-DT	PLT	.873	.931	.877	.920	.752	.885	.863	.884	.944	.865	.912*	.882
RF	ISO	.865	.934	.851	.935	.767*	<b>.920</b>	.877	.876	.933	.897*	.821	.880
BAG-DT	ISO	.867	.933	.840	.915	.749	.897	.856	.884	.940	.859	.907*	.877
SVM	PLT	.765	.886	.936	.962	.733	.866	.913*	.816	.897	.900*	.807	.862
ANN	—	.764	.884	.913	.901	.791*	.881	.932*	.859	.923	.667	.882	.854
SVM	ISO	.758	.882	.899	.954	.693*	.878	.907	.827	.897	.900*	.778	.852
ANN	PLT	.766	.872	.898	.894	.775	.871	.929*	.846	.919	.665	.871	.846
ANN	ISO	.767	.882	.821	.891	.785*	.895	.926*	.841	.915	.672	.862	.842
BST-DT	—	.874	.842	.875	.913	.523	.807	.860	.785	.933	.835	.858	.828
KNN	PLT	.819	.785	.920	.937	.626	.777	.803	.844	.827	.774	.855	.815
KNN	—	.807	.780	.912	.936	.598	.800	.801	.853	.827	.748	.852	.810
KNN	ISO	.814	.784	.879	.935	.633	.791	.794	.832	.824	.777	.833	.809
BST-STMP	PLT	.644	<b>.949</b>	.767	.688	.723	.806	.800	.862	.923	.622	.915*	.791
SVM	—	.696	.819	.731	.860	.600	.859	.788	.776	.833	.864	.763	.781
BST-STMP	ISO	.639	.941	.700	.681	.711	.807	.793	.862	.912	.632	.902*	.780
BST-STMP	—	.605	.865	.540	.615	.624	.779	.683	.799	.817	.581	.906*	.710
DT	ISO	.671	.869	.729	.760	.424	.777	.622	.815	.832	.415	.884	.709
DT	—	.652	.872	.723	.763	.449	.769	.609	.829	.831	.389	.899*	.708
DT	PLT	.661	.863	.734	.756	.416	.779	.607	.822	.826	.407	.890*	.706
LR	—	.625	.886	.195	.448	.777*	.852	.675	.849	.838	.647	.905*	.700
LR	ISO	.616	.881	.229	.440	.763*	.834	.659	.827	.833	.636	.889*	.692
LR	PLT	.610	.870	.185	.446	.738	.835	.667	.823	.832	.633	.895	.685
NB	ISO	.574	.904	.674	.557	.709	.724	.205	.687	.758	.633	.770	.654
NB	PLT	.572	.892	.648	.561	.694	.732	.213	.690	.755	.632	.756	.650
NB	—	.552	.843	.534	.556	.011	.714	-.654	.655	.759	.636	.688	.481

## Conclusion

I was fascinated to discover that from both of the studies I researched, ensemble methods were generally the best performers. I was also pleasantly surprised to find that in both studies, the top three performing algorithms were almost identical. According to Benchmarking state-of-the-art classification algorithms for credit scoring: A ten-year update, the top three classifiers were Random Forests, Bagged (MLP) Neural Networks, and Bagged Decision Trees. Meanwhile according to An Empirical Comparison of Supervised Learning Algorithms, the top three are Boosted Decision Trees, Random Forests and Bagged Decision Trees. In both studies Random Forests and Bagged Decision Trees come out on top. It's important to note that An Empirical Comparison of Supervised Learning Algorithms did not include Bagged Neural Networks. Given the agreement among both studies I have decided to use Random Forests and Bagged Decision Trees for this project. For the sake of comparison I am also going to use Support Vector Machines (SVM).

In addition to selecting classification algorithms, metrics for measuring the performance of said algorithms are required. Given the similarities of the goals of my project and the aims of Benchmarking state-of-the-art classification algorithms for credit scoring: A ten-year update, I have decided to use some of its performance measures. Specifically, Percentage Correctly Classified (PCC), and the area under a ROC curve (AUC).

## Install all required packages

```
In [1]: #install.packages("caret", repos='http://cran.us.r-project.org',dependencies =
        TRUE)
install.packages("ggplot2", repos='http://cran.us.r-project.org',dependencies
= TRUE)
install.packages("psych", repos='http://cran.us.r-project.org',dependencies =
TRUE)
install.packages("gmodels", repos='http://cran.us.r-project.org',dependencies
= TRUE)
install.packages("party", repos='http://cran.us.r-project.org',dependencies =
TRUE)
#install.packages("e1071", repos='http://cran.us.r-project.org',dependencies =
TRUE)
install.packages("unbalanced", repos='http://cran.us.r-project.org',dependenci
es = TRUE)
install.packages("hmeasure", repos='http://cran.us.r-project.org',dependencies
= TRUE)
install.packages("pROC", repos='http://cran.us.r-project.org',dependencies = T
RUE)
install.packages("adabag", repos='http://cran.us.r-project.org',dependencies =
TRUE)
library(caret)
library(psych)
library(lattice)
library(party)
library(unbalanced)
#library(hmeasure)
library(pROC)
```

package 'ggplot2' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\Ethan\AppData\Local\Temp\RtmpoD7eUN\downloaded\_packages

Warning message:

"dependencies 'graph', 'Rgraphviz' are not available"

package 'psych' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\Ethan\AppData\Local\Temp\RtmpoD7eUN\downloaded\_packages

package 'gmodels' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\Ethan\AppData\Local\Temp\RtmpoD7eUN\downloaded\_packages

package 'party' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\Ethan\AppData\Local\Temp\RtmpoD7eUN\downloaded\_packages

package 'unbalanced' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\Ethan\AppData\Local\Temp\RtmpoD7eUN\downloaded\_packages

package 'hmeasure' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\Ethan\AppData\Local\Temp\RtmpoD7eUN\downloaded\_packages

package 'pROC' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\Ethan\AppData\Local\Temp\RtmpoD7eUN\downloaded\_packages

package 'adabag' successfully unpacked and MD5 sums checked

The downloaded binary packages are in

C:\Users\Ethan\AppData\Local\Temp\RtmpoD7eUN\downloaded\_packages





```
Warning message:
"package 'caret' was built under R version 3.3.2"Loading required package: lattice
Loading required package: ggplot2
Warning message:
"package 'ggplot2' was built under R version 3.3.2"Warning message:
"package 'psych' was built under R version 3.3.2"
Attaching package: 'psych'
```

The following objects are masked from 'package:ggplot2':

%+%, alpha

```
Warning message:
"package 'party' was built under R version 3.3.2"Loading required package: grid
Loading required package: mvtnorm
Warning message:
"package 'mvtnorm' was built under R version 3.3.2"Loading required package:
modeltools
Warning message:
"package 'modeltools' was built under R version 3.3.2"Loading required package:
stats4
Loading required package: strucchange
Warning message:
"package 'strucchange' was built under R version 3.3.2"Loading required package:
zoo
```

Attaching package: 'zoo'

The following objects are masked from 'package:base':

as.Date, as.Date.numeric

```
Loading required package: sandwich
Warning message:
"package 'sandwich' was built under R version 3.3.2"Warning message:
"package 'unbalanced' was built under R version 3.3.2"Loading required package:
mlr
Warning message:
"package 'mlr' was built under R version 3.3.2"Loading required package: BBmisc
Warning message:
"package 'BBmisc' was built under R version 3.3.2"
Attaching package: 'BBmisc'
```

The following object is masked from 'package:grid':

explode

```
Loading required package: ParamHelpers
Loading required package: stringi
```

Attaching package: 'mlr'

The following object is masked `_by_` '.GlobalEnv':

db

The following object is masked from 'package:caret':

train

Loading required package: foreach

Loading required package: doParallel

Warning message:

"package 'doParallel' was built under R version 3.3.2"Loading required package: iterators

Loading required package: parallel

Warning message:

"package 'pROC' was built under R version 3.3.2"Type 'citation("pROC")' for a citation.

Attaching package: 'pROC'

The following objects are masked from 'package:stats':

cov, smooth, var

## Dataset

The dataset for this project is the Loan Data dataset from the Lending Club

(<https://www.lendingclub.com/info/download-data.action> (<https://www.lendingclub.com/info/download-data.action>)). Only 36 month term data from 2007 to February 2013 is being used because it is important that all

loans examined are well past their due date. A borrower which defaults on a current loan may do so due to a temporary job loss or a myriad of other reasons. The borrower may then recover and return to good standing once again. Using 36 month term data up until February 2013, ensures that labelling potentially temporary defaults as permanent does not occur.

### Loading Raw Data

```
In [2]: lending_club_2012_2013 <- read.csv("lending_club_rfe/LoanStats3b.csv",
      header=TRUE, skip=1)
      lending_club_2007_2011 <- read.csv("lending_club_rfe/LoanStats3a.csv",
      header=TRUE, skip=1)
```

### Pruning data set to include only 36 month term loans issued up to February 2013

The reason for this is because all 36 month loans issued February 2013 or earlier have fully come to term.

```
In [3]: lending_club_jan_2013 <- subset(lending_club_2012_2013,issue_d == "Jan-2013")
lending_club_feb_2013 <- subset(lending_club_2012_2013,issue_d == "Feb-2013")
lending_club_2012 <- subset(lending_club_2012_2013,grepl("2012",issue_d))
lending_club_2012_2013_36 <- subset(rbind(lending_club_2012,lending_club_feb_2013,lending_club_jan_2013),grepl("36",term))
lending_club_2011 <- subset(lending_club_2007_2011, grepl("2011",issue_d))
lending_club_2007_2010 <- lending_club_2007_2011[-NROW(lending_club_2011),]
lending_club_final <- rbind(lending_club_2007_2010,lending_club_2011,lending_club_2012_2013_36)
```

### Number of records in the initial dataset

Please note that each record in the dataset represents an individual loan.

```
In [4]: nrow(lending_club_final)

119276
```

### Number of features (columns) in the initial dataset

```
In [5]: ncol(lending_club_final)

111
```

### Adding label (default) to the data set

default:

1 = true (default has occurred)

0 = false (default has not occurred)

**default is the dependent variable which the predictive model aims to predict**

```
In [6]: # add label (default_status) to dataset
status <- c("Current","Fully Paid","Late (16-30 days)","Does not meet the credit policy. Status:Charged Off","Charged Off","Default","In Grace Period","Late (31-120 days)","Does not meet the credit policy. Status:Fully Paid")
default_status <- c(0,0,1,1,1,1,1,1,0)
lending_club_default <- data.frame(default = default_status[match(lending_club_final$loan_status, status)])
lending_club_final <- cbind(lending_club_final,lending_club_default)

# remove all records which have default as NA
lending_club_final <- lending_club_final[-which(is.na(lending_club_final$default)),]
```

## Cleaning data set and removing variables/features not known at issuance of loan

The dataset contains variables/features which are only known after the loan is issued. Since the aim of this project is create a predictive model to be used by lenders before the issuance of a loan, the following will be removed:

- total\_rec\_int
- total\_pymnt\_inv
- total\_pymnt
- total\_rec\_prncp
- collection\_recovery\_fee
- recoveries
- last\_pymnt\_amnt
- total\_rec\_late\_fee
- last\_pymnt\_d
- last\_pymnt\_amnt
- next\_pymnt\_d
- out\_prncp
- out\_prncp\_inv
- issue\_d
- initial\_list\_status
- funded\_amnt
- funded\_amnt\_inv
- id
- pymnt\_plan
- grade
- subgrade

```
In [7]: lending_club_final <- lending_club_final[, !(colnames(lending_club_final) %in%
c("total_rec_int","total_pymnt_inv","total_pymnt"
,"total_rec_prncp","collection_recovery_fee",
"recoveries","last_pymnt_amnt",
"total_rec_late_fee","last_pymnt_d",
"last_pymnt_amnt","next_pymnt_d","out_prncp",
"out_prncp_inv","issue_d",
"initial_list_status","funded_amnt",
"funded_amnt_inv","id","pymnt_plan","grade",
"subgrade"))]
```

## Number of Features in Dataset Before Eliminating 'All NA' Fields

```
In [8]: length(colnames(lending_club_final))
```

93

### Eliminate features which are all NA and populate NA's with column (feature) averages

```
In [9]: # find which variables/features are all NA
na_pct <- sapply(lending_club_final, function(y) sum(is.na(y))/length(y))
na_pct <- data.frame(na_pct)
all_na <- na_pct == 1

# remove variables/features which are all NA
lending_club_final <- lending_club_final[, -which(all_na==TRUE)]

# Find which variables/features have some NA values
na_pct <- sapply(lending_club_final, function(y) sum(is.na(y))/length(y))
na_pct <- data.frame(na_pct)
sig_na <- na_pct > 0
sig_na.df <- as.data.frame(sig_na)
sig_na.df <- subset(sig_na.df, sig_na.df$na_pct==TRUE)
sig_na_col <- row.names(sig_na.df)

# create a data frame of variables/features with some NA values
lending_club_final_na_sig <- lending_club_final[, which(names(lending_club_final) %in% sig_na_col)]

# remove variables/features which have some NA values
lending_club_final <- lending_club_final[, -which(sig_na==TRUE)]

# change the datatypes of the dataframe above to integer
lending_club_final_na_sig_int <- as.data.frame(lapply(lending_club_final_na_sig, as.integer))

# apply column averages to NA values
for(i in 1:ncol(lending_club_final_na_sig_int)){
  lending_club_final_na_sig_int[is.na(lending_club_final_na_sig_int[,i]), i] <-
  mean(lending_club_final_na_sig_int[,i],

      na.rm = TRUE)
}

# bring it all together
lending_club_final <- cbind(lending_club_final_na_sig_int, lending_club_final)
```

### Number of Features in Dataset After Eliminating 'All NA' Features

```
In [10]: length(colnames(lending_club_final))
```

76

## Features After Initial Cleansing

```
In [11]: # columns/features in dataset
colnames(lending_club_final)
```

```
"emp_title" "annual_inc" "title" "delinq_2yrs" "inq_last_6mths"
"mths_since_last_delinq" "mths_since_last_record" "open_acc" "pub_rec"
"total_acc" "collections_12_mths_ex_med" "mths_since_last_major_derog"
"acc_now_delinq" "tot_coll_amt" "tot_cur_bal" "total_rev_hi_lim"
"acc_open_past_24mths" "avg_cur_bal" "bc_open_to_buy" "bc_util"
"chargeoff_within_12_mths" "delinq_amnt" "mo_sin_old_il_acct"
"mo_sin_old_rev_tl_op" "mo_sin_rcnt_rev_tl_op" "mo_sin_rcnt_tl" "mort_acc"
"mths_since_recent_bc" "mths_since_recent_bc_dlq" "mths_since_recent_inq"
"mths_since_recent_revol_delinq" "num_accts_ever_120_pd" "num_actv_bc_tl"
"num_actv_rev_tl" "num_bc_sats" "num_bc_tl" "num_il_tl" "num_op_rev_tl"
"num_rev_accts" "num_rev_tl_bal_gt_0" "num_sats" "num_tl_120dpd_2m"
"num_tl_30dpd" "num_tl_90g_dpd_24m" "num_tl_op_past_12m" "pct_tl_nvr_dlq"
"percent_bc_gt_75" "pub_rec_bankruptcies" "tax_liens" "tot_hi_cred_lim"
"total_bal_ex_mort" "total_bc_limit" "total_il_high_credit_limit" "member_id"
"loan_amnt" "term" "int_rate" "installment" "sub_grade" "emp_length"
"home_ownership" "verification_status" "loan_status" "url" "desc" "purpose"
"zip_code" "addr_state" "dti" "earliest_cr_line" "revol_bal" "revol_util"
"last_credit_pull_d" "policy_code" "application_type" "default"
```

From looking at the variables, intuitively I would guess the following are significant:


1. `int_rate`: Interest rate
2. `emp_length`: Employment length in years. The longer the period of employment the more stable the income.
3. `loan_amnt`: The loan amount. The larger the loan the more difficult to pay off.
4. `annual_inc`: Annual income.
5. `delinq_2yrs`: The number of 30+ days past-due incidences of delinquency in the borrower's credit file for the past 2 years. This gives some insight as to the borrower's reliability.
6. `dti`: Debt to income ratio. The higher the ratio, the more likely the borrower is to default.

In the next section, each of the features above will be compared to default for correlation.

Please note, for a full description of all of the features in this dataset please refer to Appendix A.

### *int\_rate*

```
In [12]: library(ggplot2)
ggplot(lending_club_final, aes(as.factor(default), as.double(int_rate))) + geom_boxplot() +
labs(title="Loan Amount vs. Interest Rate", x="default status", y="interest rate")
```



### Correlation Coefficient - *int\_rate* vs. *default*

```
In [13]: # find correlation coefficient
cor(as.numeric(lending_club_final$int_rate), lending_club_final$default, method="pearson")
```

0.134073311733284

As can be seen from the boxplot above, default is more likely among higher interest rates. This is as expected. Nonetheless, given the correlation coefficient, the interest rate and default status are weakly correlated.

***emp\_length***



```
In [14]: #install.packages("gmodels", repos='http://cran.us.r-project.org',dependencies
        = TRUE)
        library(gmodels)

        emp_length.factor <- lending_club_final$emp_length
        default.factor <- as.factor(lending_club_final$default)

        emp_length.levels <- levels(emp_length.factor)
        emp_length.levels <-
        emp_length.levels[2:nlevels(lending_club_final$emp_length)]
        emp_length.len <- length(emp_length.levels)

        joint <- CrossTable(emp_length.factor, default.factor, prop.chisq=FALSE)
        joint_counts <- joint$prop.col
        barplot(joint_counts, beside=TRUE, ylab="proportion", xlab="default status", c
        ol=rainbow(emp_length.len))
        legend("topright", emp_length.levels, pch=15, col=rainbow(emp_length.len))
```

Warning message:

"package 'gmodels' was built under R version 3.3.2"

Attaching package: 'gmodels'

The following object is masked from 'package:pROC':

ci




## Cell Contents

-----
N
N / Row Total
N / Col Total
N / Table Total
-----

Total Observations in Table: 119273

emp_length.factor	default.factor		Row Total
	0	1	
< 1 year	9819	1665	11484
	0.855	0.145	0.096
	0.096	0.097	
	0.082	0.014	
1 year	7710	1267	8977
	0.859	0.141	0.075
	0.075	0.074	
	0.065	0.011	
10+ years	25862	4312	30174
	0.857	0.143	0.253
	0.253	0.252	
	0.217	0.036	
2 years	10407	1638	12045
	0.864	0.136	0.101
	0.102	0.096	
	0.087	0.014	
3 years	9087	1526	10613
	0.856	0.144	0.089
	0.089	0.089	
	0.076	0.013	
4 years	8249	1309	9558
	0.863	0.137	0.080
	0.081	0.076	
	0.069	0.011	
5 years	8592	1415	10007
	0.859	0.141	0.084
	0.084	0.083	
	0.072	0.012	
6 years	6550	1084	7634
	0.858	0.142	0.064
	0.064	0.063	
	0.055	0.009	
7 years	5133	910	6043

	0.849	0.151	0.051
	0.050	0.053	
	0.043	0.008	
-----	-----	-----	-----
8 years	4114	681	4795
	0.858	0.142	0.040
	0.040	0.040	
	0.034	0.006	
-----	-----	-----	-----
9 years	3400	524	3924
	0.866	0.134	0.033
	0.033	0.031	
	0.029	0.004	
-----	-----	-----	-----
n/a	3226	793	4019
	0.803	0.197	0.034
	0.032	0.046	
	0.027	0.007	
-----	-----	-----	-----
Column Total	102149	17124	119273
	0.856	0.144	
-----	-----	-----	-----



### Correlation Coefficient - emp\_length vs. default


```
In [15]: # find correlation coefficient
cor(as.numeric(lending_club_final$emp_length), lending_club_final$default, method="pearson")
```

0.0111569730555627

As can be seen from the diagram above, the distribution of loans by employment length is essentially the same for those who defaulted and those who did not. Therefore, emp\_length does not have much of a bearing in whether a borrower is likely or not to default. This corresponds to the fact that the correlation coefficient indicates a very weak correlation.

### Loan Amount (loan\_amnt)

```
In [16]: ggplot(lending_club_final, aes(as.factor(default), loan_amnt)) +  
  geom_boxplot() +  
  labs(title="Loan Amount vs. Default Status", x="default status", y="loan amount")
```



### Correlation Coefficient - loan amount vs. default


```
In [17]: # find correlation coefficient  
cor(as.numeric(lending_club_final$loan_amnt), lending_club_final$default, method="pearson")
```

0.0222278832849699

As can be seen from the boxplot above, the loan amount has little bearing on whether or not a borrower is going to default. This is further supported by a small correlation coefficient. Therefore we can conclude that `loan_amnt` and `default` are weakly correlated.

**Annual Income (annual\_inc)**

```
In [18]: ggplot(lending_club_final, aes(as.factor(default), annual_inc)) +
  geom_boxplot()
```



The boxplot is not visible due to extreme outliers. I will redo boxplot without outliers.

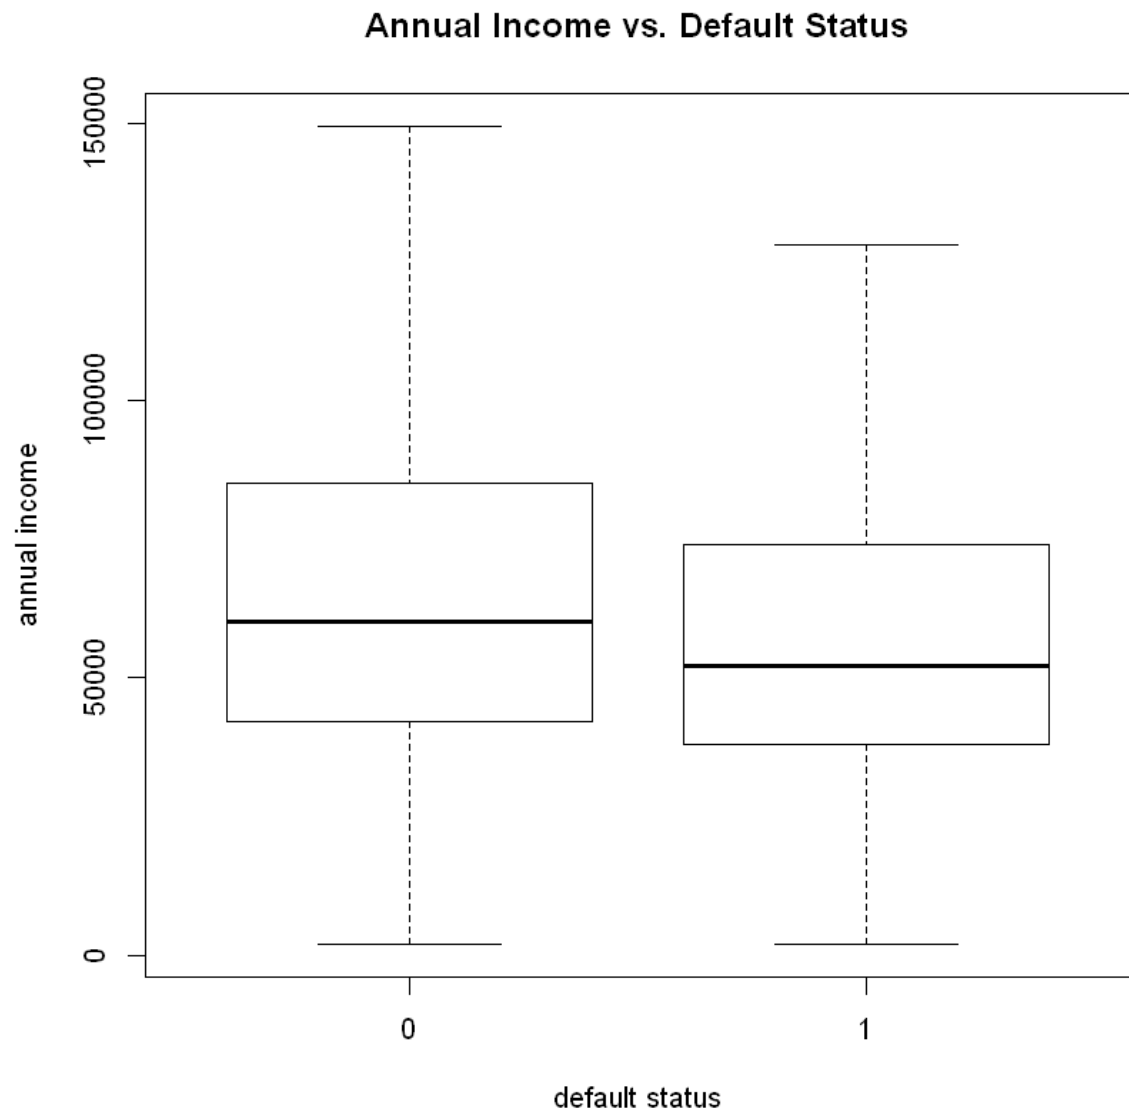
```
In [19]: summary(lending_club_final$annual_inc)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1896	41000	60000	68910	83000	7142000

The outlier problem is further illustrated by the summary table above. The minimum annual income is \$1896 and the maximum is \$7,142,000. That's a range of over 7 million!



```
In [20]: boxplot(annual_inc~as.factor(default),outline=FALSE, data=lending_club_final,
          ylab="annual income", xlab="default status")
          title("Annual Income vs. Default Status")
```



### **Correlation Coefficient - annual income vs. default status**

```
In [21]: # find correlation coefficient
          cor(as.numeric(lending_club_final$annual_inc), lending_club_final$default, method="pearson")
```

-0.0493465883580697

From the boxplot above we can see that those who defaulted were more likely to have a lower income. Nonetheless, given the correlation coefficient, the correlation is weak.

***Number of delinquencies in the last 2 years (delinq\_2yrs)***

```
In [22]: summary(lending_club_final$delinq_2yrs)
delinq_2.table <- table(delinq_2yrs=lending_club_final$delinq_2yrs,default=as.
factor(lending_club_final$default))
prop.table(delinq_2.table,2)

#install.packages("psych", repos='http://cran.us.r-project.org',dependencies =
TRUE)
#library(psych)
describeBy(lending_club_final$delinq_2yrs, as.factor(lending_club_final$default))

#library(lattice)
stripplot(lending_club_final$delinq_2yrs~as.factor(lending_club_final$default)
jitter=TRUE,
          main="Delinq. 2 Years vs. Default Status", xlab="default status", ylab="delinq_2yrs")
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0000	0.0000	0.0000	0.1726	0.0000	18.0000

```

                                default
delinq_2yrs                    0          1
0                               8.803219e-01 8.685471e-01
0.172637616987018             2.545301e-04 1.751927e-04
1                               8.841007e-02 9.530484e-02
2                               2.056799e-02 2.359262e-02
3                               6.324095e-03 7.066106e-03
4                               2.065610e-03 2.803083e-03
5                               1.047489e-03 1.167951e-03
6                               4.601122e-04 5.255781e-04
7                               2.545301e-04 3.503854e-04
8                               6.852735e-05 1.751927e-04
9                               8.810659e-05 0.000000e+00
10                              3.915848e-05 5.839757e-05
11                              4.894811e-05 1.751927e-04
12                              1.957924e-05 0.000000e+00
13                              1.957924e-05 0.000000e+00
14                              0.000000e+00 5.839757e-05
18                              9.789621e-06 0.000000e+00

```

```
$`0`
```

	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
X1	1	102149	0.17	0.56	0	0.02	0	0	18	18	5.83	61.52	0


```
$`1`
```

	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
X1	1	17124	0.19	0.61	0	0.04	0	0	14	14	5.85	60.51	0

```
attr(,"call")
```

```
by.default(data = x, INDICES = group, FUN = describe, type = type)
```

## Delinq. 2 Years vs. Default Status



### Correlation Coefficient - delinq\_2yrs vs. default

```
In [23]: # find correlation coefficient
cor(lending_club_final$delinq_2yrs, lending_club_final$default, method="pearson")
```

0.0132384016452998

The distributions for the number of delinquencies in the last 2 years by default status are very similar. Combined with a very small coefficient, we can conclude delinq\_2yrs and default are weakly correlated.

### Debt-to-Income Ratio (dti)


```
In [24]: describeBy(lending_club_final$dti, as.factor(lending_club_final$default))

ggplot(lending_club_final, aes(as.factor(default), dti)) + geom_boxplot() +
  labs(title="DTI Ratio vs. Default Status", x="default status", y="dti")

$`0`
  vars      n mean   sd median trimmed  mad min   max range skew kurtosis
X1     1 102149 14.82 7.31  14.65   14.69 8.02   0 34.99 34.99 0.16   -0.54
  se
X1 0.02

$`1`
  vars      n mean   sd median trimmed  mad min   max range skew kurtosis
X1     1 17124 15.78 7.32  15.86   15.74 7.84   0 34.95 34.95 0.05   -0.51 0.
06

attr(,"call")
by.default(data = x, INDICES = group, FUN = describe, type = type)
```



### ***Correlation Coefficient - dti vs. default***

```
In [25]: # find correlation coefficient  
cor(lending_club_final$dti, lending_club_final$default, method="pearson")  
  
0.0460055401734663
```


The boxplots above demonstrate that those who defaulted were more likely have a higher DTI (Debt-to-Income) ratio, but given that the small correlation coefficient we can conclude that dti and default are very weakly correlated.


### ***Other features considered***


I considered several other features in my analysis but they are all very heavily skewed to one value. As a result they have no predictive power. In the original dataset these features were mostly NA. When averaging was applied above, these features were converted from mostly NA to mostly the calculate average value.


```
In [26]: plot(prop.table(table(lending_club_final$delinq_amnt)), ylab="proportion", main="delinq_amnt values")
plot(prop.table(table(lending_club_final$mths_since_last_delinq)), ylab="proportion", main="mths_since_last_delinq values")
plot(prop.table(table(lending_club_final$collections_12_mths_ex_med)), ylab="proportion", main="collections_12_mths_ex_med")
plot(prop.table(table(lending_club_final$mths_since_last_major_derog)), ylab="proportion", main="mths_since_last_major_derog")
plot(prop.table(table(lending_club_final$mths_since_last_delinq)), ylab="proportion", main="mths_since_last_delinq")
```

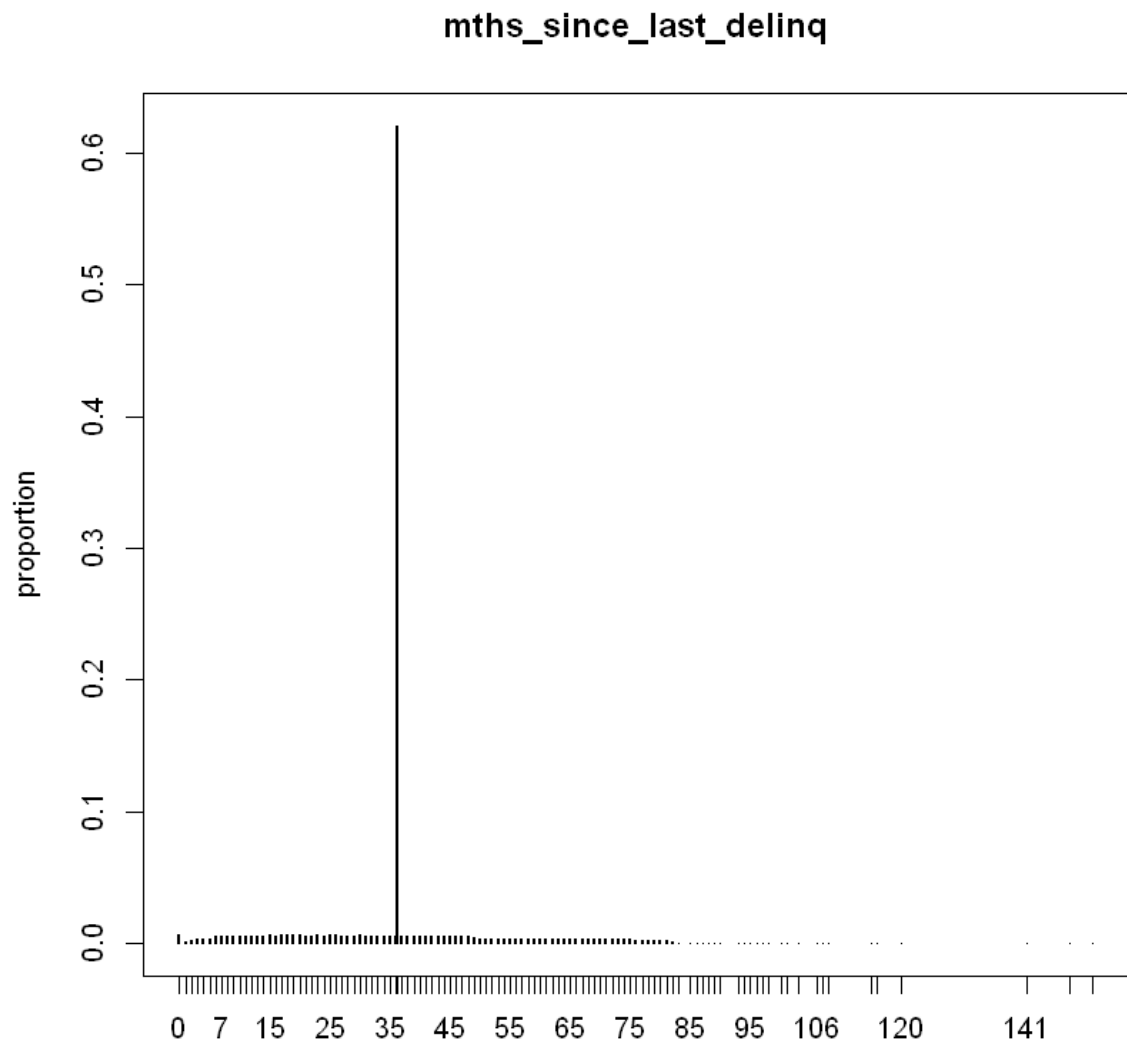


**delinq\_amnt values**

**mths\_since\_last\_delinq values**

**collections\_12\_mths\_ex\_med**





## Converting Data Types

Some data types are not accepted by various machine learning algorithms. Some of them are converted here.

```

In [27]: # convert int_rate to numeric
lending_club_final$int_rate <- as.character(lending_club_final$int_rate)
lending_club_final$int_rate <- as.numeric(substr(lending_club_final$int_rate,1
har(lending_club_final$int_rate)-1))

# convert earliest_cr_line to POSIX
lending_club_final$earliest_cr_line <- as.vector(sapply(lending_club_final$ear
liest_cr_line, function(x) paste0(x,"-01")))
lending_club_final$earliest_cr_line <- as.Date(lending_club_final$earliest_cr_
line,"%b-%Y-%d")
lending_club_final$earliest_cr_line <- as.numeric(as.POSIXct(lending_club_fina
l$earliest_cr_line, format="%Y-b%-d"))

# convert revol_util to numeric
lending_club_final$revol_util <- as.character(lending_club_final$revol_util)
lending_club_final$revol_util <- as.numeric(substr(lending_club_final$revol_ut
il,1,nchar(lending_club_final$revol_util)-1))

# convert last_credit_pull_d to POSIX because as factors they have too many le
vels!
lending_club_final$last_credit_pull_d <- as.vector(sapply(lending_club_final$la
st_credit_pull_d, function(x) paste0(x,"-01")))
lending_club_final$last_credit_pull_d <- as.Date(lending_club_final$last_credi
t_pull_d,"%b-%Y-%d")
lending_club_final$last_credit_pull_d <- as.numeric(as.POSIXct(lending_club_fi
nal$last_credit_pull_d, format="%Y-b%-d"))

# remove useless variables
to_remove <- c("url","desc","title","emp_title","id","loan_status","zip_code")
lending_club_final <- lending_club_final[ , !(names(lending_club_final) %in% t
o_remove)]

```

## Remove incomplete records

Incomplete records should be removed as long as they are a small portion of the data set < 3%

```

In [28]: nrow <- nrow(lending_club_final)
ncomplete <- sum(complete.cases(lending_club_final))
print(1-(ncomplete/nrow))

[1] 0.001224083

```

Since incomplete rows represent < 3% of data set, remove them.

## Number of records after incomplete records removed

```
In [29]: lending_club_final <- lending_club_final[complete.cases(lending_club_final),]  
  
# how many records in data set so far  
lcf_before_na_rm <- nrow(lending_club_final)  
lcf_before_na_rm  
  
119127
```

## Pair-Wise Correlations Analysis

In this section I will find pair-wise correlations. Highly correlated features are redundant and can possibly diminish the performance of a prediction model.

```
In [30]: lending_club_final_fact_as_num <- lending_club_final

# change factor columns/ features to numeric
indx <- sapply(lending_club_final_fact_as_num, is.factor)
lending_club_final_fact_as_num[indx] <-
lapply(lending_club_final_fact_as_num[indx], function(x) seq_along(levels(x)))

# create subset of dataset with only numeric features
num_feat <- sapply(lending_club_final_fact_as_num, is.numeric)
lending_club_final_num_only <- lending_club_final_fact_as_num[,num_feat]

# eliminate columns/features which are all the same value
unilength <- sapply(lending_club_final_num_only,function(x) length(unique(x)))
lending_club_final_num_only <- subset(lending_club_final_num_only, select=unilength>1)

# calculate correlation matrix
correlationMatrix <- cor(lending_club_final_num_only)

# summarize the correlation matrix
#print(correlationMatrix)


# find features that are highly correlated
highlyCorrelated <- findCorrelation(correlationMatrix, cutoff=0.6)
highCorrelationMatrix <- correlationMatrix[highlyCorrelated,highlyCorrelated]

# plot correlation of highly correlated features
library(corrplot)
corrplot(highCorrelationMatrix, method="circle")
```



Warning message:

"package 'corrplot' was built under R version 3.3.2"



## Dataset Balance

In this section I am going to determine if the dataset is balanced. In this case a balanced dataset would be one where the proportion of loans which defaulted was similar to those who did not, and vice versa for an unbalanced dataset. People normally pay off their loans so the data should be unbalanced in favour of non-default.

How many records per class? 0 = non-default, 1 = default

```
In [31]: table(lending_club_final$default)
```

```

      0      1
102039 17088
```

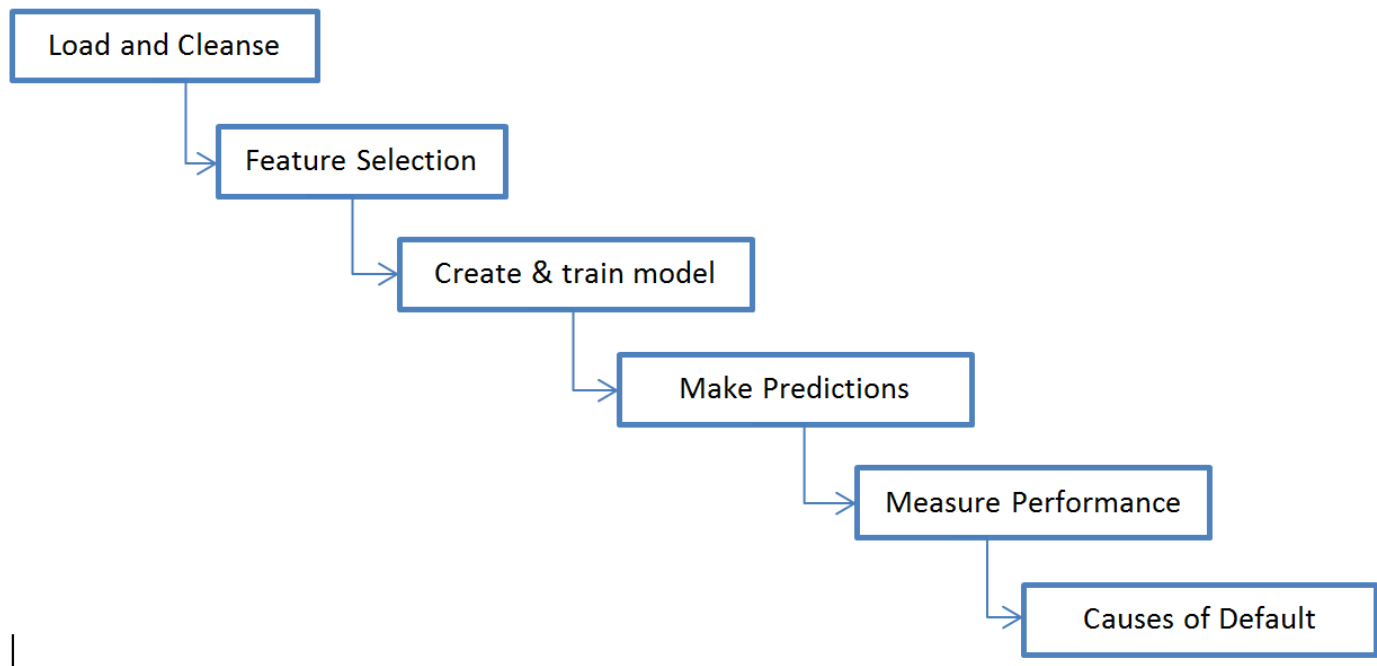
What is the proportion per class? 0 = non-default, 1 = default

```
In [32]: prop.table(table(lending_club_final$default))
```

```
      0      1  
0.8565564 0.1434436
```

As we can see the dataset is unbalanced and will have to be balanced before creating our predictive models.

# Approach



## Step 1: Load, Cleanse and Balance

Load dataset and perform initial cleansing. Cleansing will entail removing features which only contain NA values, imputing values into features which have a large proportion of NA values, changing feature data types in order to perform correlation analysis. Imputing will be done by populating NA's with the column average. Finally, the dataset will be tested for imbalance and corrected if necessary.

## Step 2: Feature Selection

Determine which features are significant and which will be used in the predictive models. Some feature analysis will be performed in order to understand the data but Recursive Feature Elimination (RFE) will be relied upon to determine the optimal feature subset.

## Step 3: Create and Train Model

Using the pruned dataset consisting of features determined from steps 1 and 2, create a test and training dataset. Create predictive models on the training dataset for the following algorithms: Random Forests, Bagged Decision Trees, and Support Vector Machines (SVM). Note: Each model is created using the same training dataset.

## Step 4: Make Predictions

Use the models on the test dataset and record the results.

## Step 5: Measure Performance

This is a sub-step of step 5. For each result set in step 5, measure the Percentage Correctly Classified (PCC), the area under a ROC curve (AUC), and accuracy, then compare the performance of each of the classification algorithms.

### Step 6: Causes of Default

Using the knowledge gained from all of the previous steps, determine the causes of loan default.

## Predictive Models

### Create training and test data sets

```
In [33]: train_rows <- sample(nrow(lending_club_final),(nrow(lending_club_final)*0.6))
lending_club.train <- lending_club_final[train_rows,]
lending_club.test <- lending_club_final[-train_rows,]
```

### Before balancing the dataset

```
In [34]: # How many records per class?
table(lending_club.train$default)

# Plot histogram of class distribution
hist(lending_club.train$default)


# Percentage of total dataset per class (0= non-default, 1= default)
prop.table(table(lending_club.train$default))

hist(lending_club.train$default)

      0      1
61143 10333

      0      1
0.855434 0.144566
```

**Histogram of lending\_club.train\$default**



### Balancing the Training Dataset - Using SMOTE

```
In [35]: #install.packages("unbalanced", repos='http://cran.us.r-project.org', dependencies = TRUE)
#library(unbalanced)

n <- ncol(lending_club.train)
y <- as.factor(lending_club.train$default)
x <- lending_club.train[, -n]
lending_club.train.smote <- ubSMOTE(X=x, Y=y)
lending_club.train.smote$default <- as.numeric(as.character(lending_club.train.smote$Y))

In [36]: lending_club.train <- cbind(lending_club.train.smote$X, default=lending_club.train.smote$default)
```

### ***After balancing the dataset***

```
In [37]: # How many records per class?
table(lending_club.train$default)

# Plot histogram of class distribution
hist(lending_club.train$default)


# Percentage of total dataset per class (0= non-default, 1= default)
prop.table(table(lending_club.train$default))

hist(lending_club.train$default)

      0      1
41332 30999

      0      1
0.5714286 0.4285714
```

**Histogram of lending\_club.train\$default**



## Random Forest

### Train model and make predictions - Pre-RFE

I am training the random forest algorithm against all of the variables/features. This model will be a baseline to compare to others.

```
In [38]: library(caret)
#install.packages("e1071", repos='http://cran.us.r-project.org',dependencies =
TRUE)
library(e1071)
library(randomForest)
set.seed(9)

ptm <- proc.time()
#####
rf.fit <- randomForest(as.factor(default)~., data=lending_club.train, importan
ce=TRUE, ntree=400)

# Generate predictions based on model
lending_club.test$default.pred <- predict(rf.fit,lending_club.test)
#####
rf.pre.time <- proc.time() - ptm
```

Warning message:

"package 'e1071' was built under R version 3.3.2"

Attaching package: 'e1071'

The following object is masked from 'package:mlr':

impute

randomForest 4.6-12

Type rfNews() to see new features/changes/bug fixes.

Attaching package: 'randomForest'

The following object is masked from 'package:psych':

outlier

The following object is masked from 'package:ggplot2':

margin

### Measure performance - Random Forest Pre-RFE



```
In [39]: # Create Confusion Matrix
cm.pre_rfe <- confusionMatrix(lending_club.test$default.pred,lending_club.test:
fault)
cm.pre_rfe

# area under a ROC curve
#auc(lending_club.test$default,as.numeric(lending_club.test$default.pred))
```

#### Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	40475	6293
1	421	462

Accuracy : 0.8591  
 95% CI : (0.8559, 0.8622)  
 No Information Rate : 0.8582  
 P-Value [Acc > NIR] : 0.2978  
  
 Kappa : 0.0912  
 McNemar's Test P-Value : <2e-16  
  
 Sensitivity : 0.98971  
 Specificity : 0.06839  
 Pos Pred Value : 0.86544  
 Neg Pred Value : 0.52322  
 Prevalence : 0.85824  
 Detection Rate : 0.84941  
 Detection Prevalence : 0.98147  
 Balanced Accuracy : 0.52905  
  
 'Positive' Class : 0

### Recursive Feature Elimination (RFE)

So far we have not considered significance when cleansing the data set. RFE analysis will identify a data set subset for optimal results.

```
In [42]: # Let's use RFE to see if we can prune some variables/features and hopefully get a better result
set.seed(77)

# 10-fold cross-validation
control <- rfeControl(functions=rfFuncs, method="cv", number=10)
lending_club.train.rfe <- lending_club.train[sample(nrow(lending_club.train),2000),]
rfe.train <- rfe(lending_club.train.rfe[,1:69], as.factor(lending_club.train.rfe[,70]), sizes=1:69, rfeControl=control)

# how big is the optimal variable subset?
print(rfe.train$bestSubset)
```

Attaching package: 'plyr'


The following object is masked from 'package:modeltools':

empty

[1] 23

## Plot the number of variables vs. RMSE

```
In [43]: plot(rfe.train, type=c("g", "o"), cex = 1.0, col = 1:69)
```



### What specific variables/features comprise the optimal subset?

```
In [44]: predictors(rfe.train)
```

"last\_credit\_pull\_d" "chargeoff\_within\_12\_mths" "collections\_12\_mths\_ex\_med"  
 "inq\_last\_6mths" "int\_rate" "mths\_since\_last\_record" "pub\_rec\_bankruptcies"  
 "num\_tl\_120dpd\_2m" "num\_tl\_90g\_dpd\_24m" "num\_tl\_30dpd" "purpose"  
 "open\_acc" "tot\_coll\_amt" "term" "annual\_inc" "dti" "loan\_amnt" "installment"  
 "num\_accts\_ever\_120\_pd" "acc\_open\_past\_24mths" "earliest\_cr\_line"  
 "member\_id" "revol\_util"

```
In [45]: predictors <- predictors(rfe.train)
formula <- paste("as.factor(default)",paste(predictors, collapse=" + "), sep="."

formula <- as.formula(formula)
formula

as.factor(default) ~ last_credit_pull_d + chargeoff_within_12_mths +
collections_12_mths_ex_med + inq_last_6mths + int_rate +
mths_since_last_record + pub_rec_bankruptcies + num_tl_120dpd_2m +
num_tl_90g_dpd_24m + num_tl_30dpd + purpose + open_acc +
tot_coll_amt + term + annual_inc + dti + loan_amnt + installment +
num_accts_ever_120_pd + acc_open_past_24mths + earliest_cr_line +
member_id + revol_util
```

### Train Random Forest model Post-RFE and generate predictions

```
In [46]: # Redo the Random Forest model with the optimal subset according to RFE
set.seed(44)

ptm <- proc.time()
#####
rf.fit.opt <- randomForest(formula, data=lending_club.train, ntree=400,
type='classification')

lending_club.test$default.pred.opt <- predict(rf.fit.opt,lending_club.test)
#####
rf.post.time <- proc.time() - ptm
```

### Measure performance - Random Forest Post-RFE

```
In [47]: # Create Confusion Matrix
cm.post_rfe <- confusionMatrix(lending_club.test$default.pred.opt,lending_club.test$default)
cm.post_rfe
```

#### Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	36590	2155
1	4306	4600

Accuracy : 0.8644  
 95% CI : (0.8613, 0.8675)  
 No Information Rate : 0.8582  
 P-Value [Acc > NIR] : 5.29e-05  
  
 Kappa : 0.5081  
 McNemar's Test P-Value : < 2.2e-16  
  
 Sensitivity : 0.8947  
 Specificity : 0.6810  
 Pos Pred Value : 0.9444  
 Neg Pred Value : 0.5165  
 Prevalence : 0.8582  
 Detection Rate : 0.7679  
 Detection Prevalence : 0.8131  
 Balanced Accuracy : 0.7878  
  
 'Positive' Class : 0

## Bagged Decision Trees

### Create model and make predictions

```
In [48]: set.seed(3)
# BAGGED DECISION TREES
library(rpart)
#install.packages("adabag", repos='http://cran.us.r-project.org',dependencies
= TRUE)
library(adabag)
lending_club.train$default.factor <- as.factor(lending_club.train$default)
lending_club.test$default.factor <- as.factor(lending_club.test$default)
formula_bdt <- paste("default.factor",paste(predictors,collapse=" +
"),sep="~")

ptm <- proc.time()
#####
# mfinal indicates total number of trees grown
# and minsplit is the minimum number of observations that must exist in a node
in order for a split to be attempted
bdt.bagging <- bagging(formula_bdt, mfinal=500, control=rpart.control(minsplit
= 50), data=lending_club.train)

# make predictions
bdt.bagging.pred <- predict.bagging(bdt.bagging, newdata=lending_club.test)
#####
bdt.time <- proc.time() - ptm

Warning message:
"package 'adabag' was built under R version 3.3.2"Loading required package: m
lbench
```

## Bagged Decision Trees - Measure Performance

```
In [49]: # Create Confusion Matrix
cm.bdt <- confusionMatrix(bdt.bagging.pred$class,lending_club.test$default)
cm.bdt

# find Area Under a ROC Curve (AUC)
#auc(lending_club.test$default,as.numeric(bdt.bagging.pred$class))
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	32953	2086
1	7943	4669

Accuracy : 0.7895  
95% CI : (0.7858, 0.7932)  
No Information Rate : 0.8582  
P-Value [Acc > NIR] : 1

Kappa : 0.3649  
McNemar's Test P-Value : <2e-16

Sensitivity : 0.8058  
Specificity : 0.6912  
Pos Pred Value : 0.9405  
Neg Pred Value : 0.3702  
Prevalence : 0.8582  
Detection Rate : 0.6915  
Detection Prevalence : 0.7353  
Balanced Accuracy : 0.7485

'Positive' Class : 0

## Support Vector Machines

```

In [50]: set.seed(10)

ptm <- proc.time()
#####
# create SVM model
lc.svm.fit <- svm(formula, data=lending_club.train)

# make predictions
lending_club.test$default.pred.svm <- predict(lc.svm.fit,lending_club.test)
#####
svm.time <- proc.time() - ptm

# Create Confusion Matrix
cm.svm <- confusionMatrix(lending_club.test$default.pred.svm,lending_club.test:
fault)
cm.svm

```

#### Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	35672	3135
1	5224	3620

Accuracy : 0.8246  
 95% CI : (0.8211, 0.828)  
 No Information Rate : 0.8582  
 P-Value [Acc > NIR] : 1  
  
 Kappa : 0.3615  
 McNemar's Test P-Value : <2e-16  
  
 Sensitivity : 0.8723  
 Specificity : 0.5359  
 Pos Pred Value : 0.9192  
 Neg Pred Value : 0.4093  
 Prevalence : 0.8582  
 Detection Rate : 0.7486  
 Detection Prevalence : 0.8144  
 Balanced Accuracy : 0.7041  
  
 'Positive' Class : 0

## Visualize the Prediction Model Using Conditional Inference Trees (CTree)



```
In [51]: lending_club.init_sol <- lending_club_final

# train and test datasets
train_rows <- sample(nrow(lending_club.init_sol), (nrow(lending_club.init_sol)*i

lending_club.init_sol.train <- lending_club.init_sol[train_rows,]
lending_club.init_sol.test <- lending_club.init_sol[-train_rows,]

lending_club.init_sol.train$default <- factor(lending_club.init_sol.train$default, levels=c(0,1), labels=c('PAID', 'DEFAULT'))
lending_club.init_sol.test$default <- factor(lending_club.init_sol.test$default, levels=c(0,1), labels=c('PAID', 'DEFAULT'))
```

```
In [52]: # create model
default_pred.ctree <- ctree(formula, data=lending_club.init_sol.train,
                           controls=ctree_control(testtype="Bonferroni", minsplit=4000, minbucket=4000))

# save model of tree to disk
png("ctree_pub.png", res=80, height=900, width=2000)
  plot(default_pred.ctree)
dev.off()
```

png: 2

## Measure Accuracy - CTree

```
In [53]: lending_club.init_sol.test$default.pred <- predict(default_pred.ctree, newdata=lending_club.init_sol.test, type='response')
cm.ctree <- confusionMatrix(lending_club.init_sol.test$default.pred, lending_club.init_sol.test$default, positive="PAID")
```

## Compile Prediction Performance Results

```
In [54]: # populate Accuracy (PCC)
acc <- cbind(as.numeric(cm.post_rfe$overall)[1],as.numeric(cm.bdt$overall)
[1],as.numeric(cm.svm$overall)[1])

# populate Kappa
kap <- cbind(as.numeric(cm.post_rfe$overall)[2],as.numeric(cm.bdt$overall)
[2],as.numeric(cm.svm$overall)[2])

# populate Sensitivity
sen <- cbind(as.numeric(cm.post_rfe$byClass)[1],as.numeric(cm.bdt$byClass)
[1],as.numeric(cm.svm$byClass)[1])

# populate Specificity
spec <- cbind(as.numeric(cm.post_rfe$byClass)[2],as.numeric(cm.bdt$byClass)[2]
numeric(cm.svm$byClass)[2])

# populate AUC
auc <- cbind(as.numeric(cm.post_rfe$byClass)[11],as.numeric(cm.bdt$byClass)[11]
s.numeric(cm.svm$byClass)[11])

perf.mat <- rbind(acc,kap,sen,spec,auc)
rownames(perf.mat) <- c("PCC (1)","Kappa (2)","Sensitivity (3)","Specificity
(4)","AUC (5)")
colnames(perf.mat) <- c("Random Forest - Post RFE","Bagged Decision Trees","SV
M")
```

## Compile Time Performance Results

```
In [55]: rf.mat <- as.matrix(rf.post.time)
svm.mat <- as.matrix(svm.time)
bdt.mat <- as.matrix(bdt.time)

et <- cbind(rf.mat[3], svm.mat[3], bdt.mat[3])
colnames(et) <- c("Random Forests","SVM","Bagged DT")
```

# Results

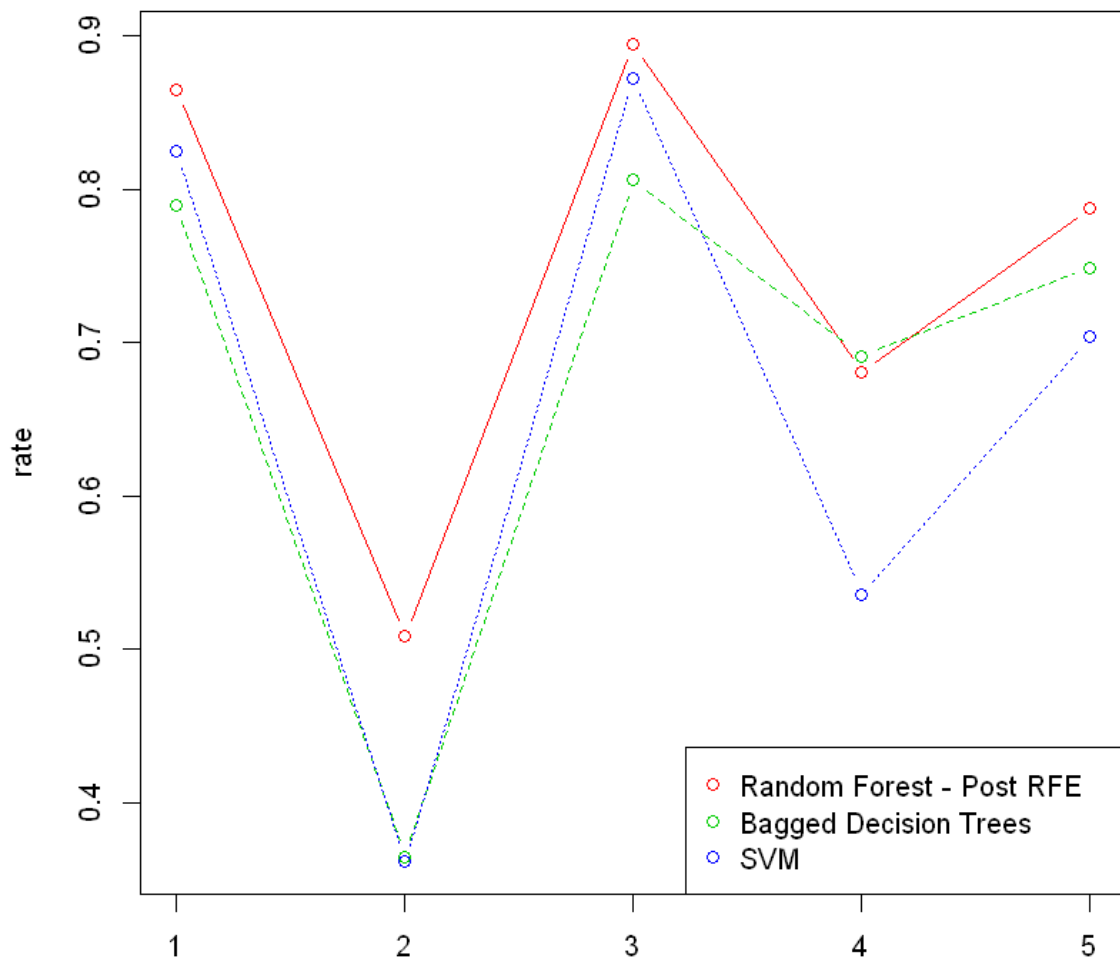
## Predictive Models Performance

This section will summarize the results of the prediction performance measures for all of the classification algorithms covered in this project. The measures of performance to be compared are Accuracy (Percent Correctly Classified), Area Under a ROC Curve (AUC), as well as Sensitivity and Specificity.

```
In [56]: matplotlib(perf.mat, type = c("b"),pch=1,col = 2:4, ylab="rate", main="Predictive  
Models Performance Results")  
legend("bottomright", legend = c("Random Forest - Post RFE","Bagged Decision T  
rees","SVM"), col=2:4, pch=1)  
perf.mat
```

	Random Forest - Post RFE	Bagged Decision Trees	SVM
<b>PCC (1)</b>	0.8644100	0.7895322	0.8245787
<b>Kappa (2)</b>	0.5081445	0.3649014	0.3614966
<b>Sensitivity (3)</b>	0.8947085	0.8057756	0.8722613
<b>Specificity (4)</b>	0.6809771	0.6911917	0.5358993
<b>AUC (5)</b>	0.7878428	0.7484837	0.7040803

**Predictive Models Performance Results**



## Predictive Models Results Summary

Over all of the metrics used, the best performer is Random Forests. Random Forest only came in second once to Bagged Decision Trees when it came to specificity. In regards to overall performance of Bagged Decision Trees and SVM it was a mixed bag. According to Benchmarking state-of-the-art classification algorithms for credit scoring: A ten-year update<sup>2</sup>, which compared Random Forests, Bagged Decision Trees and SVM among many others, Random Forests and Bagged Decision Trees outperformed SVM. The results here agree with the study in regards to Random Forests but the performance of Bagged Decision Trees was surprisingly lacking.

### *Description of Metrics Used*


The first metric is Percentage Correctly Classified (PCC), which measures the classified loans as a percentage of all loans. The second is Kappa, which is a measure of agreement between two raters. This measure is considered more robust than PCC and other similar metrics<sup>4</sup>. The third is Sensitivity which is the proportion of borrowers who were predicted as not defaulting versus all the borrowers who actually did not default. In other words it measures the success rate of predicting who will not default. The fourth is Specificity which is the proportion of borrowers who were predicted as defaulting versus all the borrowers who did default. This metric is especially important since it captures the success rate of predicting who will default. The final metric is Area Under a ROC Curve which measures the relationship between the true positive rate (TPR) and the false positive rate (FPR). An AUC of 1 means that all borrowers predicted not default did not default in reality. An AUC of 0.5 would indicate that of the borrowers predicted not to default, half actually did while the other half did not.

## Time Performance

This section will summarize the time performance for each predictive model.

```
In [57]: et
barplot(et, main="Classification Model Time Performance", ylab="elapsed time i
n seconds")
```

Random Forests	SVM	Bagged DT
112.25	1286.65	2410.80



### Time Performance Summary

The time performance results show that Random Forests outperform Support Vector Machines and Bagged Decision Trees. Again, the best time performance comes from Random Forests meanwhile Bagged Decision Trees algorithm unacceptably slow, especially for systems that require quick turnaround such as targeted display advertising.

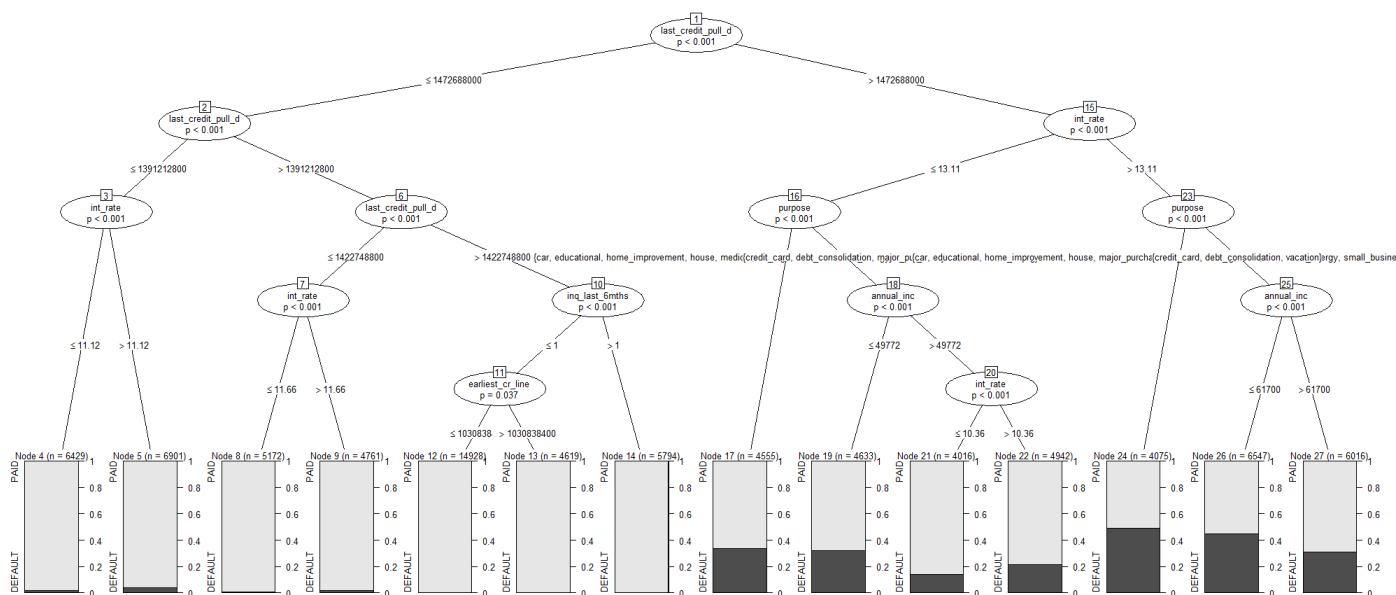
## Important Features for Prediction

This section lists the features in the dataset which are the most important for predicting loan default as determined by Recursive Feature Elimination. These features should be collected for each potential borrower to accurately determine his/her likelihood of default.

```
In [58]: dataDict <- read.table("LCDataDictionary.csv", header = TRUE, sep=",")
dataDict[which(dataDict$Feature %in% predictors(rfe.train)),]
```

	Feature	Description
2	acc_open_past_24mths	Number of trades opened in past 24 months.
5	annual_inc	The self-reported annual income provided by the borrower during registration.
11	chargeoff_within_12_mths	Number of charge-offs within 12 months
13	collections_12_mths_ex_med	Number of collections in 12 months excluding medical collections
28	inq_last_6mths	The number of inquiries in past 6 months (excluding auto and mortgage inquiries)
29	installment	The monthly payment owed by the borrower if the loan originates.
30	int_rate	Interest Rate on the loan
32	last_credit_pull_d	The most recent month LC pulled credit for this loan
37	loan_amnt	The listed amount of the loan applied for by the borrower. If at some point in time the credit department reduces the loan amount then it will be reflected in this value.
40	member_id	A unique LC assigned Id for the borrower member.
59	pub_rec_bankruptcies	Number of public record bankruptcies
60	purpose	A category provided by the borrower for the loan request.
64	revol_util	Revolving line utilization rate or the amount of credit the borrower is using relative to all available revolving credit.
67	term	The number of payments on the loan. Values are in months and can be either 36 or 60.
69	tot_coll_amt	Total collection amounts ever owed

## Loan Approval Decision Process



This decision tree above can be used to predict whether a potential borrower will or will not default on their loan. By starting at the top node of the tree and by following the decisions you eventually arrive at a probability of default or non-default (paid). By studying this tree we can glean some important insights:

The last credit pull date (`last_credit_pull_d`) is the most indicative feature when sizing up a potential borrower. If a potential borrower has had a credit pull within 4 years and 7 months ago or earlier this is of significant concern. In fact, if the borrower has not had a credit pull in the last 4 years and 7 months, the risk for default is minimal.

Another very important feature is interest rate. If the borrower's interest rate is greater than 13.11% and the he or she has had a credit pull in the last 4 years and 7 months, the probability of default ranges from 30% to 50%.

Another important feature is the purpose of the loan. If the reason for borrowing is not for credit card, debt consolidation, or vacation, the lender should take note.

**NOTES: This tree is based on the Conditional Inference Trees model generated above (see heading "Visualize the Prediction Model Using Conditional Inference Trees (CTree)"). Its performance measures are displayed below.**

This tree has been compressed and does not show all of the possible decision points. This has been done for display purposes. A larger more detailed tree can be generated for a more in-depth analysis.

\*Please note that 147000960 translates to Jan. 27, 2012 and since the data is current to Sept. 1, 2016, it roughly translates to 4 years and 7 months ago.

In [63]: cm.ctree

## Confusion Matrix and Statistics

	Reference	
Prediction	PAID	DEFAULT
PAID	30677	5062
DEFAULT	0	0

Accuracy : 0.8584

95% CI : (0.8547, 0.862)

No Information Rate : 0.8584

P-Value [Acc &gt; NIR] : 0.5037

Kappa : 0

McNemar's Test P-Value : &lt;2e-16

Sensitivity : 1.0000

Specificity : 0.0000

Pos Pred Value : 0.8584

Neg Pred Value : NaN

Prevalence : 0.8584

Detection Rate : 0.8584

Detection Prevalence : 1.0000

Balanced Accuracy : 0.5000

'Positive' Class : PAID

## Conclusion

The purpose of this project is to determine the characteristics or features of a borrower which are important in determining their creditworthiness and create a framework by which a lender could examine those features and come to a conclusion as to whether a potential borrower is or is not likely to default. The characteristics/features provided above help a lender determine what kind of data he or she needs to compile in order to screen potential borrowers. These features are listed under the **Important Features for Prediction** section above. Once the lender has gathered all of the data he or she will need a framework to turn the raw data into a decision. The framework provided in this project is the decision tree provided in the **Loan Approval Decision Process** section above. By following the decision tree and asking each question at each node in regards to a potential borrower, the lender will eventually arrive at general probability of default. With this information, he or she can decide on whether or not to issue the loan. If the borrower decides to lend, the potential rate of default can help determine the level of risk and the appropriate rate of interest to compensate. In addition to creating a decision process for each potential new borrower, the decision tree provides general insights which can be used for making quick decisions.

Overall the best performing classification model was Random Forests. With an accuracy rate of 86%, it can be useful in making predictions. In predicting non-default the model is pretty strong, where it is weaker is in predicting default.



## Real Life Example

In this section I am going to take the important characteristics/features from a few individual borrowers in the Lending Club dataset and make a predictions using the Random Forest-Post RFE model.

```
In [64]: lc_tiny_test <- lending_club.test[sample(nrow(lending_club.test),15),1:70]

lc_tiny_test$default.pred <- predict(rf.fit.opt,lc_tiny_test)

lc_tiny_test <- cbind(lc_tiny_test[,which(colnames(lc_tiny_test) %in% predictors(rfe.train))],
                      default=lc_tiny_test[,70],
                      default.pred=lc_tiny_test[,71])
lc_tiny_test
```

	annual_inc	inq_last_6mths	mths_since_last_record	open_acc	collections_12_m
<b>212310</b>	105000	2	74.666935483871	10	0
<b>126220</b>	120000	0	74.666935483871	7	0
<b>33713</b>	64500	3	74.666935483871	11	0
<b>120858</b>	89000	0	74.666935483871	10	0
<b>20168</b>	43000	1	74.666935483871	7	0
<b>179642</b>	78000	0	74.666935483871	8	0
<b>90421</b>	67000	1	74.666935483871	9	0
<b>107751</b>	55000	2	74.666935483871	10	0
<b>47791</b>	85000	2	74.666935483871	12	0
<b>144503</b>	51700	3	74.666935483871	17	0
<b>170143</b>	24000	0	74.666935483871	6	0
<b>38612</b>	115275	0	74.666935483871	4	0
<b>132395</b>	34500	1	50	8	0
<b>8253</b>	47000	0	74.666935483871	4	0
<b>3837</b>	72000	1	74.666935483871	9	0

## Appendix A - Data Dictionary

In [65]: dataDict

	Feature	Description
1	acc_now_delinq	The number of accounts on which the borrower is now delinquent.
2	acc_open_past_24mths	Number of trades opened in past 24 months.
3	addr_state	The state provided by the borrower in the loan application
4	all_util	Balance to credit limit on all trades
5	annual_inc	The self-reported annual income provided by the borrower during registration.
6	annual_inc_joint	The combined self-reported annual income provided by the co-borrowers during registration
7	application_type	Indicates whether the loan is an individual application or a joint application with two co-borrowers
8	avg_cur_bal	Average current balance of all accounts
9	bc_open_to_buy	Total open to buy on revolving bankcards.
10	bc_util	Ratio of total current balance to high credit/credit limit for all bankcard accounts.
11	chargeoff_within_12_mths	Number of charge-offs within 12 months
12	collection_recovery_fee	post charge off collection fee
13	collections_12_mths_ex_med	Number of collections in 12 months excluding medical collections
14	delinq_2yrs	The number of 30+ days past-due incidences of delinquency in the borrowers credit file for the past 2 years delinq_amnt,The past-due amount owed for the accounts on which the borrower is now delinquent. desc,Loan description provided by the borrower dti,A ratio calculated using the borrower?s total monthly debt payments on the total debt obligations excluding mortgage and the requested LC loan divided by the borrower?s self-reported monthly income. dti_joint,A ratio calculated using the co-borrowers total monthly payments on the total debt obligations excluding mortgages and the requested LC loan divided by the co-borrowers combined self-reported monthly income earliest_cr_line,The month the borrowers earliest reported credit line was opened
15	emp_length	Employment length in years. Possible values are between 0 and 10 where 0 means less than one year and 10 means ten or more years.
16	emp_title	The job title supplied by the Borrower when applying for the loan.*

	Feature	Description
17	fico_range_high	The upper boundary range the borrower?s FICO at loan origination belongs to.
18	fico_range_low	The lower boundary range the borrower?s FICO at loan origination belongs to.
19	funded_amnt	The total amount committed to that loan at that point in time.
20	funded_amnt_inv	The total amount committed by investors for that loan at that point in time.
21	grade	LC assigned loan grade
22	home_ownership	The home ownership status provided by the borrower during registration. Our values are: RENT OWN MORTGAGE OTHER.
23	id	A unique LC assigned ID for the loan listing.
24	il_util	Ratio of total current balance to high credit/credit limit on all install acct
25	initial_list_status	The initial listing status of the loan. Possible values are ? W F
26	inq_fi	Number of personal finance inquiries
27	inq_last_12m	Number of credit inquiries in past 12 months
28	inq_last_6mths	The number of inquiries in past 6 months (excluding auto and mortgage inquiries)
29	installment	The monthly payment owed by the borrower if the loan originates.
30	int_rate	Interest Rate on the loan
...	...	...
44	mo_sin_rcnt_tl	Months since most recent account opened
45	mort_acc	Number of mortgage accounts.

	Feature	Description
46	mths_since_last_delinq	<p>The number of months since the borrowers last delinquency. mths_since_last_major_derog, Months since most recent 90-day or worse rating</p> <p>mths_since_last_record, The number of months since the last public record. mths_since_rcnt_il, Months since most recent installment accounts opened</p> <p>mths_since_recent_bc, Months since most recent bankcard account opened. mths_since_recent_bc_dlq, Months since most recent bankcard delinquency</p> <p>mths_since_recent_inq, Months since most recent inquiry. mths_since_recent_revol_delinq, Months since most recent revolving delinquency. next_pymnt_d, Next scheduled payment date num_accts_ever_120_pd, Number of accounts ever 120 or more days past due</p> <p>num_actv_bc_tl, Number of currently active bankcard accounts num_actv_rev_tl, Number of currently active revolving trades num_bc_sats, Number of satisfactory bankcard accounts num_bc_tl, Number of bankcard accounts num_il_tl, Number of installment accounts num_op_rev_tl, Number of open revolving accounts num_rev_accts, Number of revolving accounts num_rev_tl_bal_gt_0, Number of revolving trades with balance &gt;0 num_sats, Number of satisfactory accounts num_tl_120dpd_2m, Number of accounts currently 120 days past due (updated in past 2 months)</p> <p>num_tl_30dpd, Number of accounts currently 30 days past due (updated in past 2 months)</p> <p>num_tl_90g_dpd_24m, Number of accounts 90 or more days past due in last 24 months</p> <p>num_tl_op_past_12m, Number of accounts opened in past 12 months open_acc, The number of open credit lines in the borrowers credit file.</p>
47	open_acc_6m	Number of open trades in last 6 months
48	open_il_12m	Number of installment accounts opened in past 12 months
49	open_il_24m	Number of installment accounts opened in past 24 months
50	open_il_6m	Number of currently active installment trades
51	open_rv_12m	Number of revolving trades opened in past 12 months
52	open_rv_24m	Number of revolving trades opened in past 24 months
53	out_prncp	Remaining outstanding principal for total amount funded
54	out_prncp_inv	Remaining outstanding principal for portion of total amount funded by investors
55	pct_tl_nvr_dlq	Percent of trades never delinquent

	Feature	Description
56	percent_bc_gt_75	Percentage of all bankcard accounts > 75% of limit.
57	policy_code	publicly available policy_code=1 new products not publicly available policy_code=2
58	pub_rec	Number of derogatory public records
59	pub_rec_bankruptcies	Number of public record bankruptcies
60	purpose	A category provided by the borrower for the loan request.
61	pymnt_plan	Indicates if a payment plan has been put in place for the loan
62	recoveries	post charge off gross recovery
63	revol_bal	Total credit revolving balance
64	revol_util	Revolving line utilization rate or the amount of credit the borrower is using relative to all available revolving credit.
65	sub_grade	LC assigned loan subgrade
66	tax_liens	Number of tax liens
67	term	The number of payments on the loan. Values are in months and can be either 36 or 60.
68	title	The loan title provided by the borrower
69	tot_coll_amt	Total collection amounts ever owed
70	tot_cur_bal	Total current balance of all accounts
71	tot_hi_cred_lim	Total high credit/credit limit
72	total_acc	The total number of credit lines currently in the borrowers credit file total_bal_ex_mort, Total credit balance excluding mortgage total_bal_il, Total current balance of all installment accounts total_bc_limit, Total bankcard high credit/credit limit total_cu_tl, Number of finance trades total_il_high_credit_limit, Total installment high credit/credit limit total_pymnt, Payments received to date for total amount funded total_pymnt_inv, Payments received to date for portion of total amount funded by investors total_rec_int, Interest received to date total_rec_late_fee, Late fees received to date total_rec_prncp, Principal received to date total_rev_hi_lim ÿ, Total revolving high credit/credit limit url, URL for the LC page with listing data. verification_status, Indicates if income was verified by LC not verified or if the income source was verified verified_status_joint, Indicates if the co-borrowers joint income was verified by LC not verified or if the income source was verified



	Feature	Description
73	zip_code	The first 3 numbers of the zip code provided by the borrower in the loan application.

## References

1. Baesens, B., Van Gestel, T., Viaene, S., Stepanova, M., Suykens, J., & Vanthienen, J. (2003). Benchmarking state-of-the-art classification algorithms for credit scoring. Journal of the Operational Research Society, 54(6), 627-635.
2. Lessmann, S., Seow, H., Baesens, B. and Thomas, L. (2013). Benchmarking state-of-the-art classification algorithms for credit scoring: A ten-year update. 1st ed. [ebook] pp.1-32. Available at: [http://www.business-school.ed.ac.uk/waf/crc\\_archive/2013/42.pdf](http://www.business-school.ed.ac.uk/waf/crc_archive/2013/42.pdf) ([http://www.business-school.ed.ac.uk/waf/crc\\_archive/2013/42.pdf](http://www.business-school.ed.ac.uk/waf/crc_archive/2013/42.pdf)) [Accessed 15 Sep. 2016].
3. Caruana, R. and Niculescu-Mizil, A. (2006). An Empirical Comparison of Supervised Learning Algorithms. 1st ed. [ebook] Pittsburgh, Pennsylvania, USA: Cornell University, pp.1-7. Available at: <https://www.cs.cornell.edu/~caruana/ctp/ct.papers/caruana.icml06.pdf> (<https://www.cs.cornell.edu/~caruana/ctp/ct.papers/caruana.icml06.pdf>) [Accessed 16 Sep. 2016].
4. En.wikipedia.org. (2016). Cohen's kappa. [online] Available at: [https://en.wikipedia.org/wiki/Cohen's\\_kappa](https://en.wikipedia.org/wiki/Cohen's_kappa) ([https://en.wikipedia.org/wiki/Cohen's\\_kappa](https://en.wikipedia.org/wiki/Cohen's_kappa)) [Accessed 24 Nov. 2016].

In [ ]:





